

01/15/24

A Step Back

There are different layers of memory going from the CPU to storage

- Registers: few, incredibly fast, very expensive
- L1 cache: bigger, slower
- L2 cache: bigger, slower
- RAM: bigger, slower
 - Speed in nanoseconds
- Secondary storage: SDD/HDD
 - Speed in milliseconds
 - Lots of storage
 - Persistent
 - Incredibly slow
 - Writing/reading data in blocks, which can vary in number of bytes
 - Even if you only need to read a smaller amount of data, you have to read the block size

Database systems: minimize HDD/SDD accesses

Sorted array of 128 integers (8 bytes each)

- Worst case binary search on 128 integers is way faster than a single additional disk access
- In a binary tree, if the number of values stored in each node increases by one, then the number of (possible) children increases by one
 - Maximizing the number of values in each node means less disk access → faster

- Minimize the height of the tree

B+ Tree

